

CHAPTER 3



Introduction to SQL

Practice Exercises

- 3.1 Write the following queries in SQL, using the university schema. (We suggest you actually run these queries on a database, using the sample data that we provide on the web site of the book, *db-book.com*. Instructions for setting up a database, and loading sample data, are provided on the above web site.)
- Find the titles of courses in the Comp. Sci. department that have 3 credits.
 - Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.
 - Find the highest salary of any instructor.
 - Find all instructors earning the highest salary (there may be more than one with the same salary).
 - Find the enrollment of each section that was offered in Fall 2017.
 - Find the maximum enrollment, across all sections, in Fall 2017.
 - Find the sections that had the maximum enrollment in Fall 2017.

Answer:

- a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

```
select title
from   course
where  dept_name = 'Comp. Sci.' and credits = 3
```

- b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.
This query can be answered in several different ways. One way is as follows.

```

select distinct t takes.ID
from takes, instructor, teaches
where takes.course_id = teaches.course_id and
      takes.section_id = teaches.section_id and
      takes.semester = teaches.semester and
      takes.year = teaches.year and
      teaches.instructor_id = instructor.id and
      instructor.name = 'Einstein'

```

- c. Find the highest salary of any instructor.

```

select max(salary)
from instructor

```

- d. Find all instructors earning the highest salary (there may be more than one with the same salary).

```

select ID, name
from instructor
where salary = (select max(salary) from instructor)

```

- e. Find the enrollment of each section that was offered in Fall 2017.

```

select course_id, section_id,
       (select count(ID)
        from takes
        where takes.year = section.year
              and takes.semester = section.semester
              and takes.course_id = section.course_id
              and takes.section_id = section.section_id)
       as enrollment
from section
where semester = 'Fall'
and year = 2017

```

Note that if the result of the subquery is empty, the aggregate function `count` returns a value of 0.

One way of writing the query might appear to be:

```

select takes. course_id, takes.section_id, count(id)
from section, takes
where takes. course_id = section. course_id
and takes.section_id = section.section_id
and takes.semester = section.semester
and takes.year = section.year
and takes.semester = 'Fall'
and takes.year = 2017
group by takes. course_id, takes.section_id

```

But note that if a section does not have any students taking it, it would not appear in the result. One way of ensuring such a section appears with a count of 0 is to use the outer join operation, covered in Chapter 4.

- f. Find the maximum enrollment, across all sections, in Fall 2017. One way of writing this query is as follows:

```

select max(enrollment)
from (select count(id) as enrollment
from section, takes
where takes.year = section.year
and takes.semester = section.semester
and takes. course_id = section. course_id
and takes.section_id = section.section_id
and takes.semester = 'Fall'
and takes.year = 2017
group by takes. course_id, takes.section_id)

```

As an alternative to using a nested subquery in the from clause, it is possible to use a with clause, as illustrated in the answer to the next part of this question.

A subtle issue in the above query is that if no section had any enrollment, the answer would be empty, not 0. We can use the alternative using a subquery, from the previous part of this question, to ensure the count is 0 in this case.

- g. Find the sections that had the maximum enrollment in Fall 2017. The following answer uses a with clause, simplifying the query.

```

with se_enrollment as (
    select takes.course_id, takes.se_id, count(id) as enrollment
    from section, takes
    where takes.year = section.year
          and takes.semester = section.semester
          and takes.course_id = section.course_id
          and takes.se_id = section.se_id
          and takes.semester = 'Fall'
          and takes.year = 2017
    group by takes.course_id, takes.se_id)
select course_id, se_id
from se_enrollment
where enrollment = (select max(enrollment) from se_enrollment)

```

It is also possible to write the query without the `with` clause, but the subquery to find enrollment would get repeated twice in the query.

While not in order to add distinct in the count, it is not necessary in light of the primary key constraint on takes.

- 3.2 Suppose you are given a relation `grade_points(grade, points)` that provides a conversion from letter grades in the takes relation to numerical scores; for example, an `A` grade could be specified to correspond to 4 points, an `A-` to 3.7 points, a `B+` to 3.3 points, a `B` to 3 points, and so on. The grade points earned by a student for a course offering (section) is defined as the number of credits for the course multiplied by the numerical points for the grade that the student received.

Given the preceding relation, and our university's schema, write each of the following queries in SQL. You may assume for simplicity that no takes tuple has the null value for grade.

- Find the total grade points earned by the student with ID 12345, across all courses taken by the student.
- Find the grade point average (GPA) for the above student, that is, the total grade points divided by the total credits for the associated courses.
- Find the ID and the grade-point average of each student.
- Now reconsider your answers to the earlier parts of this exercise under the assumption that some grades might be null. Explain whether your solutions still work and, if not, provide versions that handle nulls properly.

Answer:

- Find the total grade-points earned by the student with ID 12345, across all courses taken by the student.

```

select sum(redits * points)
from takes, course, grade_points
where takes.grade = grade_points.grade
      and takes.course_id = course.course_id
      and id = 12345

```

In the above query, a student who has not taken any course would not have any tuples, whereas we would expect to get 0 as the answer. One way of fixing this problem is to use the outer join operation, which we study later in Chapter 4. Another way to ensure that we get 0 as the answer is via the following query:

```

(select sum(redits * points)
 from takes, course, grade_points
 where takes.grade = grade_points.grade
       and takes.course_id = course.course_id
       and id = 12345 )
union
(select 0
 from student
 where id = 12345 and
       not exists ( select * from takes where id = 12345 ))

```

- b. Find the grade point average (GPA) for the above student, that is, the total grade-points divided by the total credits for the associated courses.

```

select sum(redits * points)/sum(redits) as GPA
from takes, course, grade_points
where takes.grade = grade_points.grade
      and takes.course_id = course.course_id
      and id = 12345

```

As before, a student who has not taken any course would not appear in the above result; we can ensure that such a student appears in the result by using the modified query from the previous part of this question. However, an additional issue in this case is that the sum of credits would also be 0, resulting in a divide-by-zero condition. In fact, the only meaningful way of defining the GPA in this case is to define it as null. We can ensure that such a student appears in the result with a null GPA by adding the following union clause to the above query.

```

union
(select null as GPA
 from student
 where id = 12345 and
       not exists ( select * from takes where id = 12345 ))

```

- Find the ID and the grade-point average of each student.

```

select t ID, sum(redits * points)/sum(redits) as GPA
from takes, course, grade_points
where takes.grade = grade_points.grade
      and takes.course_id = course.course_id
group by ID

```

Again, to handle students who have not taken any course, we would have to add the following union clause:

```

union
(select t ID, null as GPA
 from student
 where not exists ( select t * from takes where takes.ID = student.ID))

```

- d. Now reconsider your answers to the earlier parts of this exercise under the assumption that some grades might be null. Explain whether your solutions still work and, if not, provide versions that handle nulls properly. The queries listed above all include a test of equality on grade between grade_points and takes. Thus, for any takes tuple with a null grade, that student's course would be eliminated from the rest of the computation of the result. As a result, the credits of such courses would be eliminated also, and thus the queries would return the correct answer even if some grades are null.

3.3 Write the following inserts, deletes, or updates in SQL, using the university schema.

- a. Increase the salary of each instructor in the Comp. Sci. department by 10%.
- b. Delete all courses that have never been offered (i.e., do not occur in the section relation).
- c. Insert every student whose totalred attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.

Answer:

- a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

```

update instructor
set salary = salary * 1.10
where dept_name = Comp. Sci.

```

- b. Delete all courses that have never been offered (that is, do not occur in the section relation).

```

person (driver_id, name, address)
ar (license_plate, model, year)
a ident (report_number, year, location)
owns (driver_id, license_plate)
participated (report_number, license_plate, driver_id, damage_amount)

```

Figure 3.17 Insurance database

```

delete from course
where course_id not in
(select course_id from section)

```

- Insert every student whose total red attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.

```

insert into instructor
select id, name, dept_name, 10000
from student
where total_red > 100

```

3.4 Consider the insurance database of Figure 3.17, where the primary keys are underlined. Construct the following SQL queries for this relational database.

- a. Find the total number of people who owned cars that were involved in accidents in 2017.
- b. Delete all year-2010 cars belonging to the person whose ID is 12345.

Answer:

- a. Find the total number of people who owned cars that were involved in accidents in 2017.

Note: This is not the same as the total number of accidents in 2017. We must count people with several accidents only once. Furthermore, note that the question asks for owners, and it might be that the owner of the car was not the driver actually involved in the accident.

```

select count(distinct person.driver_id)
from accident, participated, person, owns
where accident.report_number = participated.report_number
and owns.driver_id = person.driver_id
and owns.license_plate = participated.license_plate
and year = 2017

```

- b. Delete all year-2010 cars belonging to the person whose ID is 12345.

```
delete car
where year = 2010 and license_plate in
(select license_plate
from owns o
where o.driver_id = 12345 )
```

Note: The owns, accident and participated records associated with the deleted cars still exist.

- 3.5 Suppose that we have a relation *marks*(ID, score) and we wish to assign grades to students based on the score as follows: grade F if score < 40, grade C if 40 ≤ score < 60, grade B if 60 ≤ score < 80, and grade A if 80 ≤ score. Write SQL queries to do the following:

- a. Display the grade for each student, based on the marks relation.
b. Find the number of students with each grade.

Answer:

- a. Display the grade for each student, based on the marks relation.

```
select ID,
       case
         when score < 40 then 'F'
         when score < 60 then 'C'
         when score < 80 then 'B'
         else 'A'
       end
from marks
```

- b. Find the number of students with each grade.


```

with grades as
(
    select id,
           case
               when score < 40 then 'F'
               when score < 60 then 'C'
               when score < 80 then 'B'
               else 'A'
           end as grade
    from marks
)
select grade, count(id)
from grades
group by grade
    
```

As an alternative, the `with` clause can be removed, and instead the definition of `grades` can be made a subquery of the main query.

- 3.6 The SQL-like operator `case` is case-sensitive (in most systems), but the `lower()` function on strings can be used to perform case-insensitive matching. To show how, write a query that finds departments whose names contain the string `si` as a substring, regardless of the case.

Answer:

```

select dept_name
from department
where lower(dept_name) like '%si%'
    
```

- 3.7 Consider the SQL query

```

select p.a1
from p, r1, r2
where p.a1 = r1.a1 or p.a1 = r2.a1
    
```

Under what conditions does the preceding query select values of `p.a1` that are either in `r1` or in `r2`? Examine carefully the cases where either `r1` or `r2` may be empty.

Answer:

The query selects those values of `p.a1` that are equal to some value of `r1.a1` or `r2.a1` if and only if both `r1` and `r2` are non-empty. If one or both of `r1` and `r2` are empty, the Cartesian product of `p`, `r1` and `r2` is empty, hence the result of the query is empty. If `p` itself is empty, the result is empty.

- 3.8 Consider the bank database of Figure 3.18, where the primary keys are underlined. Construct the following SQL queries for this relational database.

```

branch(branch_name, branch_city, assets)
customer(ID, customer_name, customer_street, customer_city)
loan(loan_number, branch_name, amount)
borrower(ID, loan_number)
account(account_number, branch_name, balance)
depositor(ID, account_number)

```

Figure 3.18 Banking database.

- a. Find the ID of each customer of the bank who has an account but not a loan.
- b. Find the ID of each customer who lives on the same street and in the same city as customer 12345.
- c. Find the name of each branch that has at least one customer who has an account in the bank and who lives in Harrison.

Answer:

- a. Find the ID of each customer of the bank who has an account but not a loan.

```

(select ID
from depositor)
except
(select ID
from borrower)

```

- b. Find the ID of each customer who lives on the same street and in the same city as customer 12345.

```

select F.ID
from customer as F, customer as S
where F.customer_street = S.customer_street
and F.customer_city = S.customer_city
and S.customer_id = 12345

```

- c. Find the name of each branch that has at least one customer who has an account in the bank and who lives in Harrison.

```

select distinct branch_name
from account, depositor, customer
where customer.id = depositor.id
      and depositor.account_number = account.account_number
      and customer.city = 'Harrison'

```

3.9 Consider the relational database of Figure 3.19, where the primary keys are underlined. Give an expression in SQL for each of the following queries.

- Find the ID, name, and city of residence of each employee who works for First Bank Corporation .
- Find the ID, name, and city of residence of each employee who works for First Bank Corporation and earns more than \$10000.
- Find the ID of each employee who does not work for First Bank Corporation .
- Find the ID of each employee who earns more than every employee of Small Bank Corporation .
- Assume that companies may be located in several cities. Find the name of each company that is located in every city in which Small Bank Corporation is located.
- Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).
- Find the name of each company whose employees earn a higher salary, on average, than the average salary at First Bank Corporation .

Answer:

- Find the ID, name, and city of residence of each employee who works for First Bank Corporation .

```

employee ( ID, person_name, street, city)
works ( ID, company_name, salary)
company ( company_name, city)
manages ( ID, manager_id)

```

Figure 3.19 Employee database.

```

select e.ID, e.person_name, ity
from employee as e, works as w
where w.company_name = First Bank Corporation and
      w.ID = e.ID

```

- b. Find the ID, name, and city of residence of each employee who works for First Bank Corporation and earns more than \$10000.

```

select t *
from employee
where ID in
  (select ID
   from works
   where company_name = First Bank Corporation and salary > 10000)

```

This could be written also in the style of the answer to part a.

- c. Find the ID of each employee who does not work for First Bank Corporation.

```

select ID
from works
where company_name <> First Bank Corporation

```

If one allows people to appear in employee without appearing also in works, the solution is slightly more complicated. An outer join as discussed in Chapter 4 could be used as well.

```

select ID
from employee
where ID not in
  (select ID
   from works
   where company_name = First Bank Corporation )

```

- d. Find the ID of each employee who earns more than every employee of Small Bank Corporation.

```

select ID
from works
where salary > all
  (select salary
   from works
   where company_name = Small Bank Corporation )

```

If people may work for several companies and we wish to consider the total earnings of each person, the problem is more complex. But note that the

fact that *ID* is the primary key for *works* implies that this cannot be the case.

- e. Assume that companies may be located in several cities. Find the name of each company that is located in every city in which Small Bank Corporation is located.

```
select S. company_name
from company as S
where not exists ((select ity
                  from company
                  where company_name = Small Bank Corporation )
except
(select ity
 from company as T
 where S. company_name = T. company_name))
```

- f. Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).

```
select company_name
from works
group by company_name
having count (distinct ID) >= all
(select count (distinct ID)
 from works
 group by company_name)
```

- g. Find the name of each company whose employees earn a higher salary, on average, than the average salary at First Bank Corporation .

```
select company_name
from works
group by company_name
having avg (salary) > (select avg (salary)
                      from works
                      where company_name = First Bank Corporation )
```

- 3.10 Consider the relational database of Figure 3.19. Give an expression in SQL for each of the following:

- Modify the database so that the employee whose *ID* is 12345 now lives in Newtown .
- Give each manager of First Bank Corporation a 10 per cent raise unless the salary becomes greater than \$100000; in such cases, give only a 3 per cent raise.

Answer:

- a. Modify the database so that the employee whose ID is 12345 now lives in Newtown .

```
update employee
set city = Newtown
where id = 12345
```

- b. Give each manager of First Bank Corporation a 10 per cent raise unless the salary becomes greater than \$100000; in such cases, give only a 3 per cent raise.

```
update works T
set T.salary = T.salary * 1.03
where T.id in (select manager_id
               from manages)
and T.salary * 1.1 > 100000
and T.company_name = First Bank Corporation
```

```
update works T
set T.salary = T.salary * 1.1
where T.id in (select manager_id
               from manages)
and T.salary * 1.1 <= 100000
and T.company_name = First Bank Corporation
```

The above updates would give different results if executed in the opposite order. We give below a safer solution using the case statement.

```
update works T
set T.salary = T.salary <
( case
  when (T.salary < 1.1 > 100000) then 1.03
  else 1.1
end)
where T.id in (select manager_id
               from manages) and
T.company_name = First Bank Corporation
```